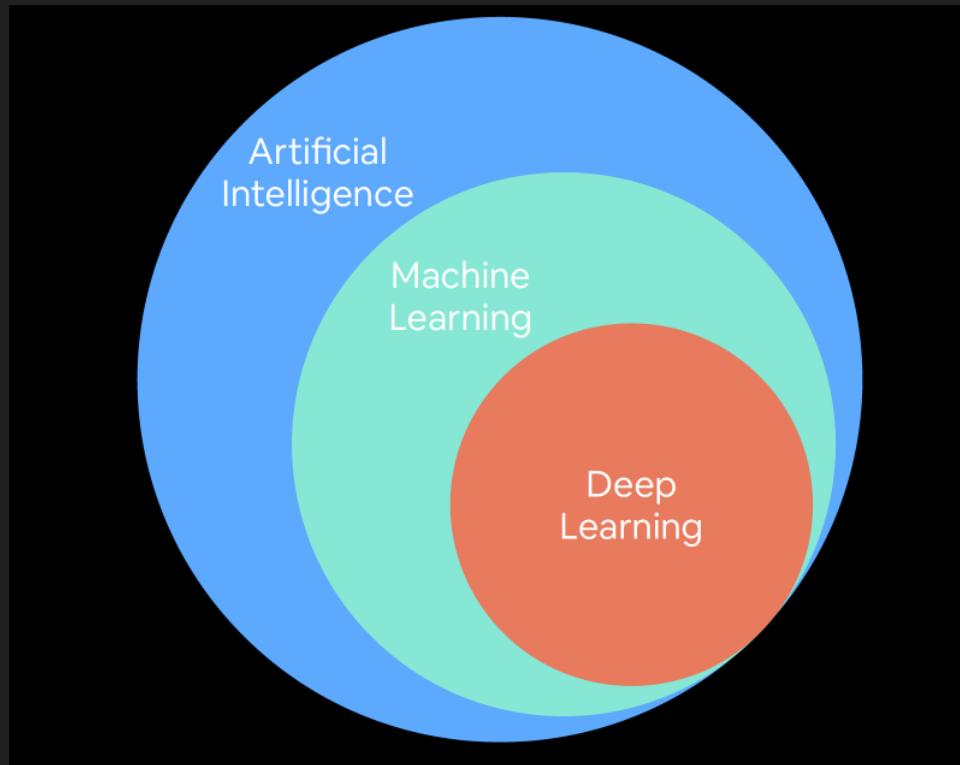


โครงการพัฒนาทักษะการเรียนรู้ของเครื่อง (Machine Learning) ของบัณฑิตเพื่อตอบสนองการพัฒนาประเทศไทย 4.0

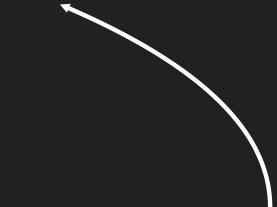
16 และ 23 กุมภาพันธ์ 2566



Machine Learning vs. Deep Learning



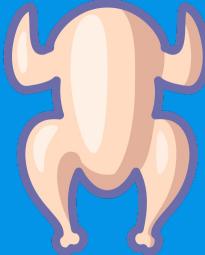
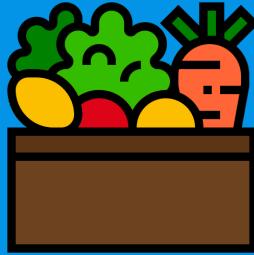
Machine learning is turning things (data) into numbers and **finding patterns** in those numbers.



The computer does this part.
How?
Code & math.
We're going to be writing the code.

Traditional programming

Inputs



Rules

1. Cut vegetables
2. Season chicken
3. Preheat oven
4. Cook chicken for 30-minutes
5. Add vegetables



Output

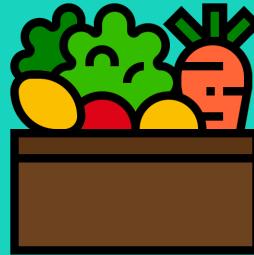


Starts with

Makes

Machine learning algorithm

Inputs



Output



Rules

1. Cut vegetables
2. Season chicken
3. Preheat oven
4. Cook chicken for 30-minutes
5. Add vegetables

Starts with

Figures out

Good reason: Why not?

Better reason: For a complex problem, can you think of all the rules?

(probably not)

(maybe not very simple...)

“If you can build a **simple rule-based system that doesn’t require machine learning, do that.”**

— A wise software engineer... (actually rule 1 of [Google’s Machine Learning Handbook](#))

What deep learning is good for



- **Problems with long lists of rules**—when the traditional approach fails, machine learning/deep learning may help.
- **Continually changing environments**—deep learning can adapt ('learn') to new scenarios.
- **Discovering insights within large collections of data**—can you imagine trying to hand-craft rules for what 101 different kinds of food look like?

(typically)

What deep learning is not good for



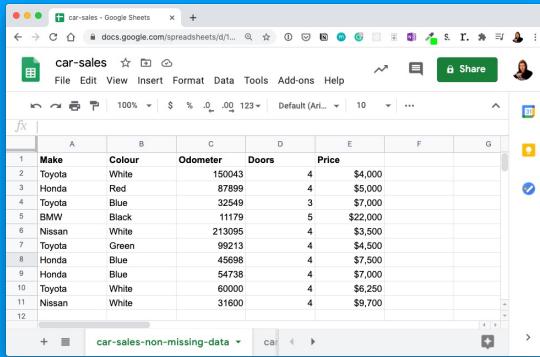
- **When you need explainability**—the patterns learned by a deep learning model are typically uninterpretable by a human.
- **When the traditional approach is a better option** — if you can accomplish what you need with a simple rule-based system.
- **When errors are unacceptable** — since the outputs of deep learning model aren't always predictable.
- **When you don't have much data** — deep learning models usually require a fairly large amount of data to produce great results.

(though we'll see how to get great results without huge amounts of data)

Machine Learning

Algorithm: gradient boosted machine

dmlc
XGBoost



Make	Colour	Odometer	Doors	Price
Toyota	White	150043	4	\$4,000
Honda	Red	87899	4	\$5,000
Toyota	Blue	32549	3	\$7,000
BMW	Black	11179	5	\$22,000
Nissan	White	213095	4	\$3,500
Toyota	Green	99213	4	\$4,500
Honda	Blue	45698	4	\$7,500
Honda	Blue	54738	4	\$7,000
Toyota	White	60000	4	\$6,250
Nissan	White	31600	4	\$9,700



Structured data

Deep Learning



Daniel Bourke @mrdbourke · Nov 1
"How do I learn #machinelearning?"

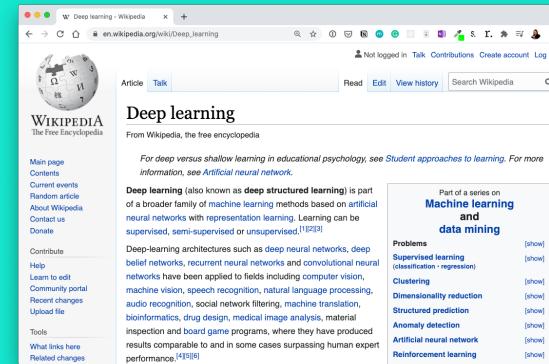
What you want to hear:

1. Learn Python
2. Learn Math/Stats/Probability
3. Learn software engineering
4. Build

What you need to do:

1. Google it
2. Go down the rabbit hole
3. Resurface in 6-9 months and reassess

See you on the other side.



Deep learning

From Wikipedia, the free encyclopedia

For deep versus shallow learning in educational psychology, see [Student approaches to learning](#). For more information, see [Artificial neural network](#).

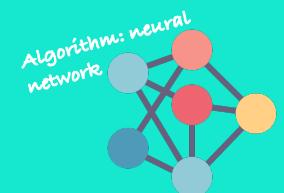
Deep learning (also known as [deep structured learning](#)) is part of a broader family of machine learning methods based on [artificial neural networks](#) with [representation learning](#). Learning can be supervised, semi-supervised or unsupervised.^{[1][2]}

Deep-learning architectures such as [deep neural networks](#), [deep belief networks](#), [recurrent neural networks](#) and [convolutional neural networks](#) have been applied to fields including [computer vision](#), [machine vision](#), [speech recognition](#), [natural language processing](#), [audio recognition](#), [social network filtering](#), [machine translation](#), [biometrics](#), [drug design](#), [medical image analysis](#), [material inspection](#) and [game programs](#), where they have produced results comparable to and in some cases surpassing human expert performance.^{[3][4]}

Part of a series on
Machine learning and
data mining

- Problems
 - [Supervised learning](#) [show]
 - [Classification - regression](#) [show]
 - [Clustering](#) [show]
 - [Dimensionality reduction](#) [show]
 - [Structured prediction](#) [show]
 - [Anomaly detection](#) [show]
 - [Artificial neural network](#) [show]
 - [Reinforcement learning](#) [show]

Unstructured data



Machine Learning vs. Deep Learning

(common algorithms)

- Random forest
- Gradient boosted models
- Naive Bayes
- Nearest neighbour
- Support vector machine
- ...many more

(since the advent of deep learning these are
often referred to as “shallow algorithms”)

- Neural networks
- Fully connected neural network
- Convolutional neural network
- Recurrent neural network
- Transformer
- ...many more

What we're focused on building
(with PyTorch)

(depending how you represent your
problem, many algorithms can be
used for both)

Structured data

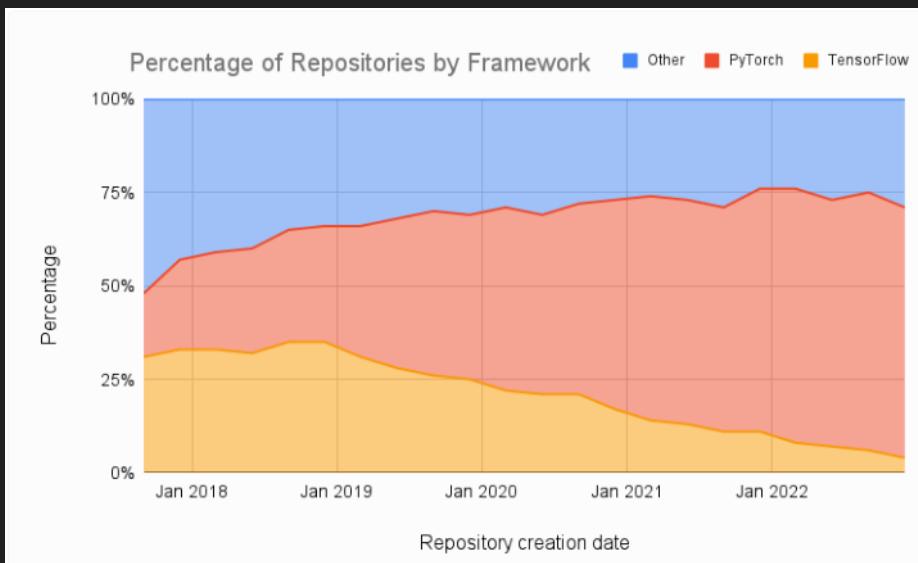
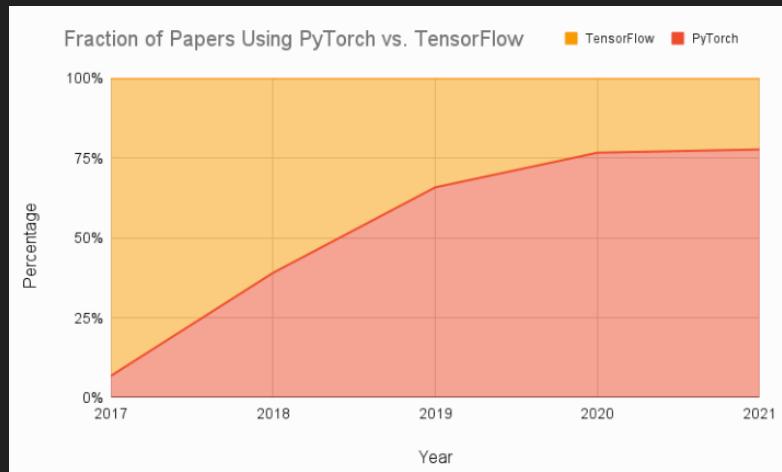
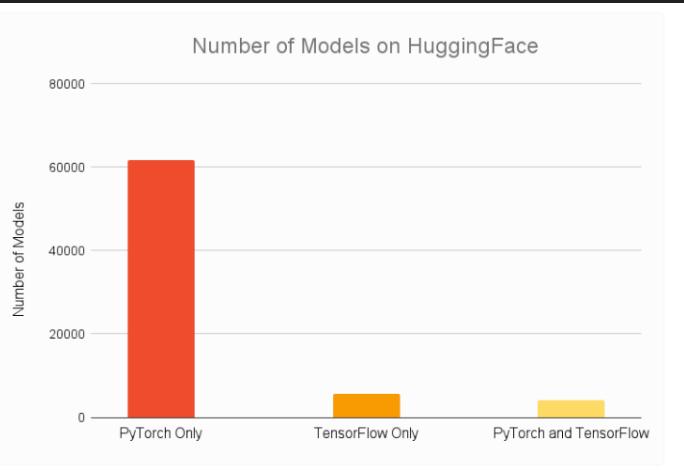
Unstructured data



What is PyTorch?



- Most popular research deep learning framework*
- Write fast deep learning code in Python (able to run on a GPU/many GPUs)
- Able to access many pre-built deep learning models (Torch Hub/
torchvision.models)
- Whole stack: preprocess data, model data, deploy model in your application/cloud
- Originally designed and used in-house by Facebook/Meta (now open-source and used by companies such as Tesla, Microsoft, OpenAI)



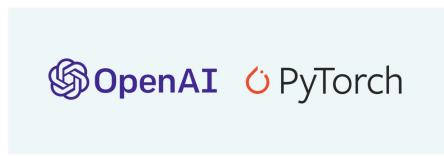
<https://www.assemblyai.com/blog/pytorch-vs-tensorflow-in-2023/>



A screenshot of a YouTube video player. The video shows a man, Andrej Karpathy, standing on a stage at a Tesla event. He is pointing towards a large screen behind him which displays the text "PLEASE WELCOME ANDREJ KARPATHY SR DIRECTOR OF AI AT TESLA". The video player interface includes a progress bar at 0:25 / 11:10, chapters, and various sharing options. The title of the video is "PyTorch at Tesla - Andrej Karpathy, Tesla".

OpenAI Standardizes on PyTorch

We are standardizing OpenAI's deep learning framework on PyTorch. In the past, we implemented projects in many frameworks depending on their relative strengths. We've now chosen to standardize to make it easier for our team to create and share optimized implementations of our models.

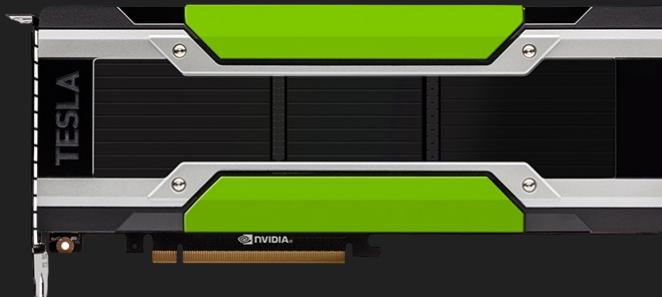


A screenshot of a GitHub repository page for "PyTorch". The repository has 8400+ stars and 1700+ forks. It is described as "The Incredible PYTORCH". The page lists a curated list of tutorials, projects, libraries, videos, papers, books and anything related to PyTorch. Below the repository details, there is a news article titled "AI for AG: Production machine learning for agriculture" by Chris Padwick, Director of Computer Vision and Machine Learning at Blue River Technology. The article discusses how farming affects daily life in cities and how PyTorch is used in agriculture. The date of the article is Aug 7, 2020, and it has a 11 min read time.

A screenshot of a blog post from Meta AI titled "PyTorch builds the future of AI and machine learning at Facebook". The post is dated June 2, 2021. It discusses how Facebook's AI models perform trillions of inference operations per day and how PyTorch is used to handle this workload. The post includes sharing options for Facebook and Twitter, and a "Our Work" section with a small image of a field. The date of the post is January 30, 2020, and it has a 1 minute read time.



Microsoft



GPU (Graphics Processing Unit)



TPU (Tensor Processing Unit)

Neural Networks



(before data gets used with a neural network, it needs to be turned into numbers)

[[116, 78, 15],
 [117, 43, 96],
 [125, 87, 23],
 ...,

Daniel Bourke @mrbourke - Nov 1
"How do I learn machinelearning?"

What you want to hear:
1. Learn Python
2. Learn Math/Stats/Probability
3. Learn software engineering
4. Build

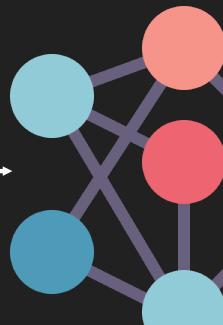
What you need to do:
1. Google it
2. Go down the rabbit hole
3. Resurface in 6-9 months and reassess

See you on the other side.



Inputs

Each of these nodes is called a "hidden unit" or "neuron".



(choose the appropriate neural network for your problem)

Learns representation (patterns/features/weights)

Representation outputs

[[0.983, 0.004, 0.013],
 [0.110, 0.889, 0.001],
 [0.023, 0.027, 0.985],
 ...,

(a human can understand these)

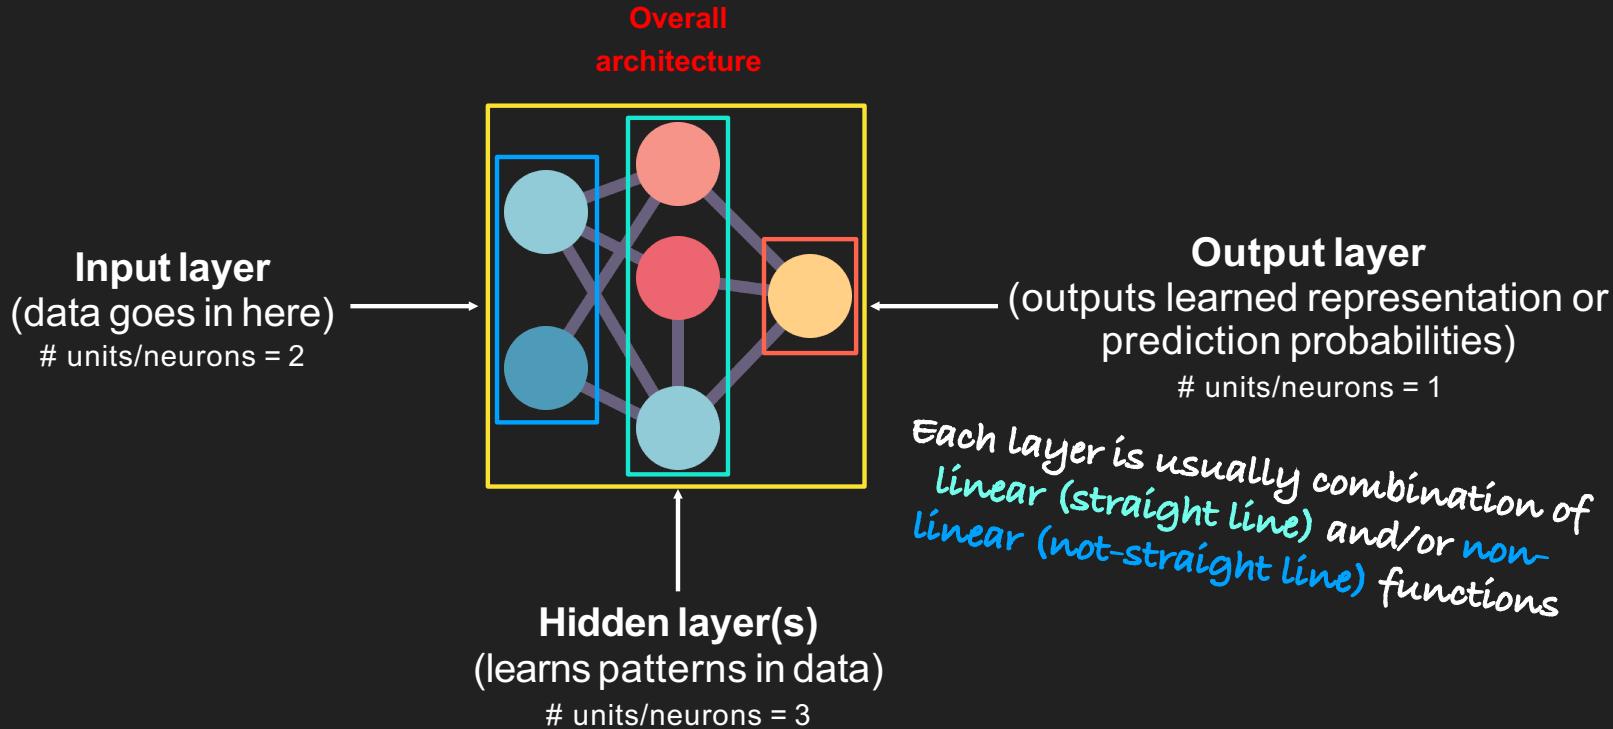
Ramen,
Spaghetti

Not a disaster

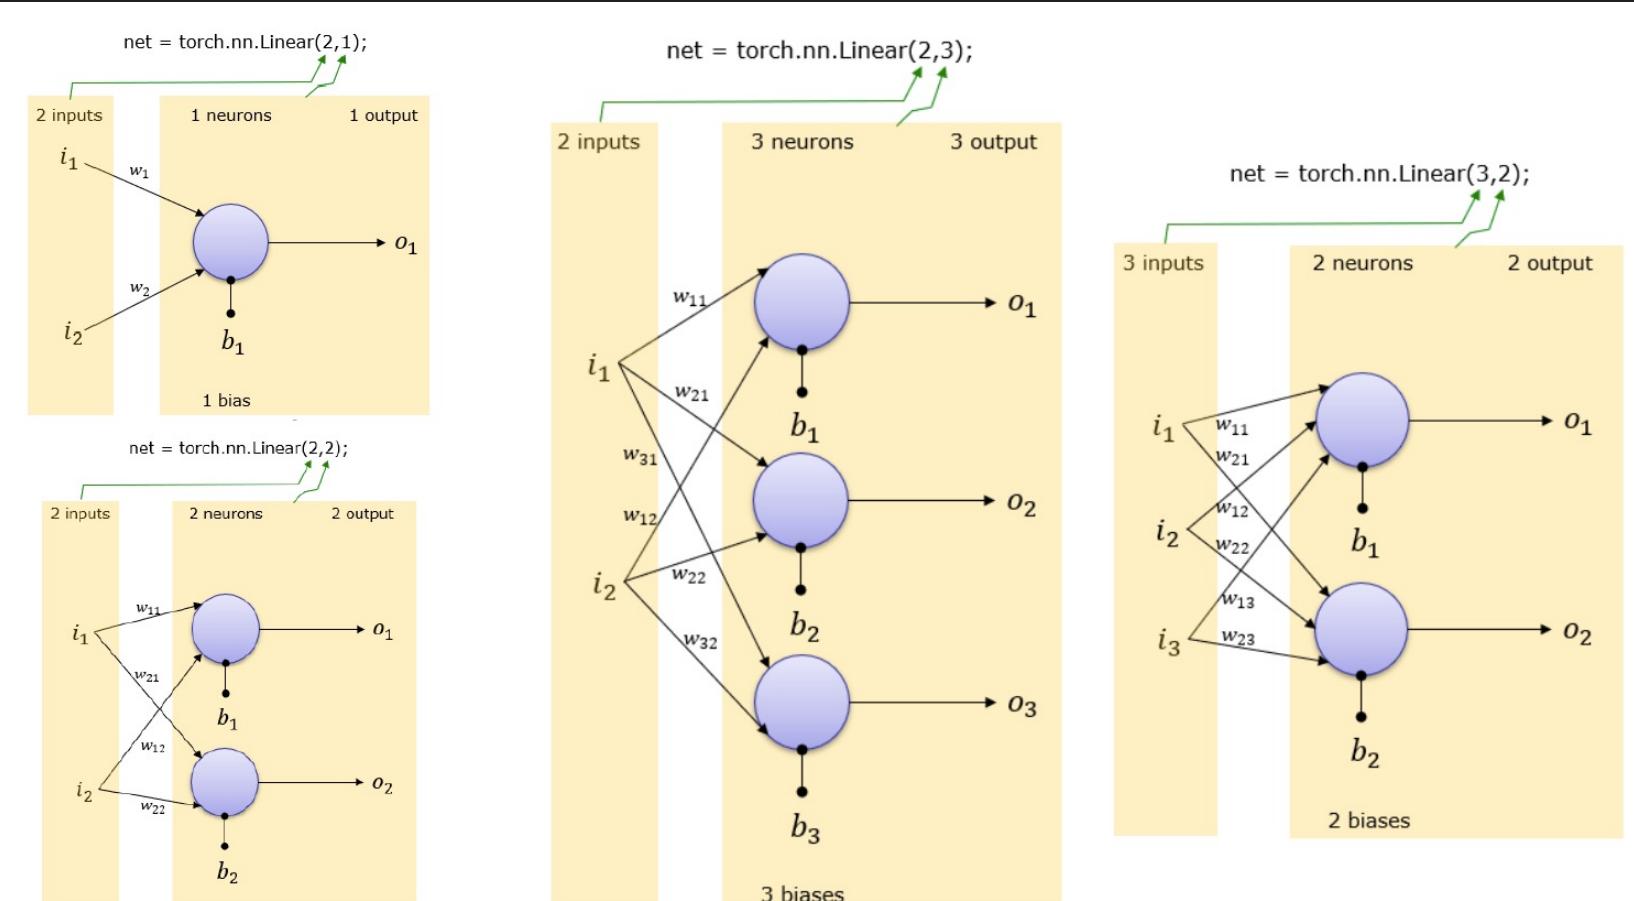
"Hey Siri, what's the weather today?"

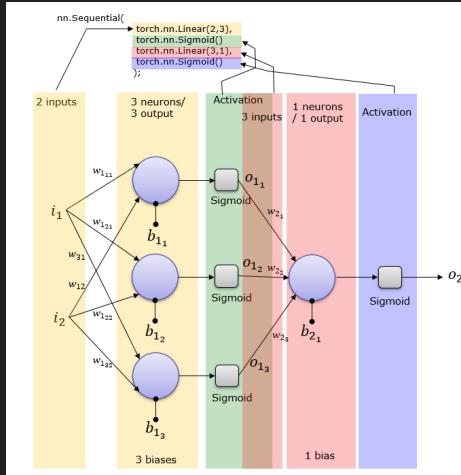
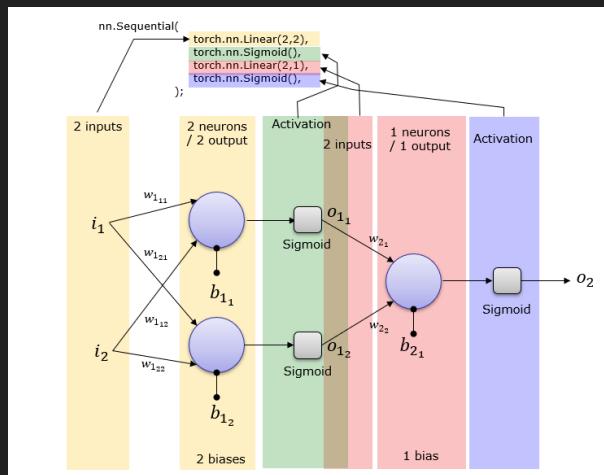
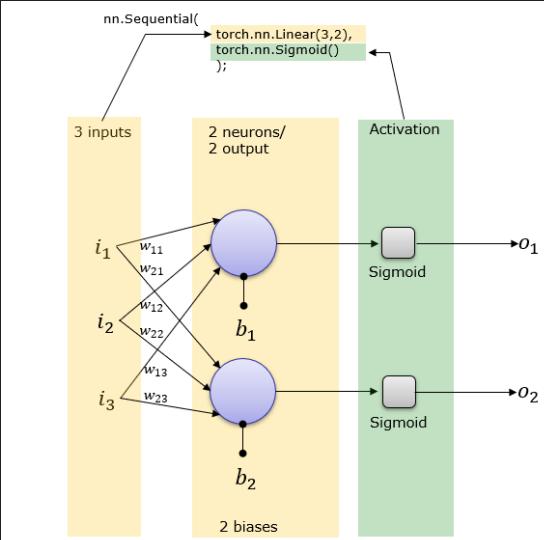
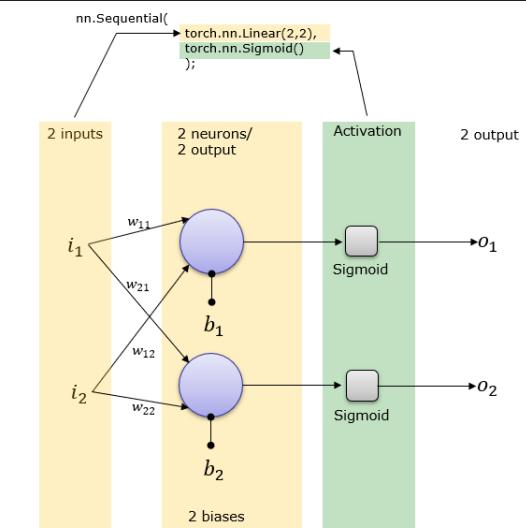
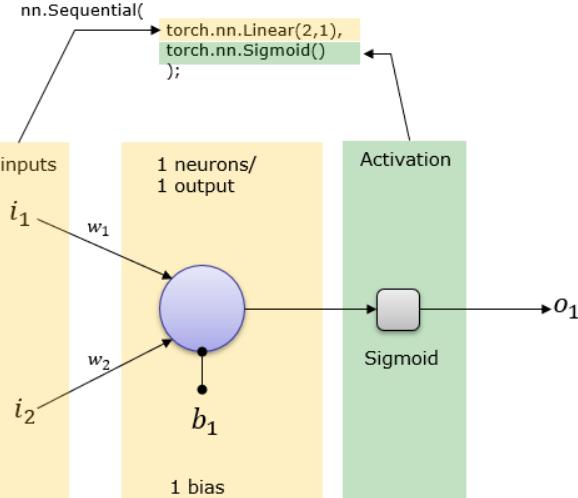
Outputs

Anatomy of Neural Networks



Note: “patterns” is an arbitrary term, you’ll often hear “embedding”, “weights”, “feature representation”, “feature vectors” all referring to similar things.





Three Ways to Build a Neural Network in PyTorch

```
from torch import nn
from collections import OrderedDict

# define model architecture
model3 = nn.Sequential(OrderedDict([
    ('fc1', nn.Linear(16, 12)),
    ('relu1', nn.ReLU()),
    ('fc2', nn.Linear(12, 10)),
    ('relu2', nn.ReLU()),
    ('fc3', nn.Linear(10, 1)),
    ('sigmoid', nn.Sigmoid())
]))

# print model architecture
print(model3)
```

```
Sequential(
  (fc1): Linear(in_features=16, out_features=12, bias=True)
  (relu1): ReLU()
  (fc2): Linear(in_features=12, out_features=10, bias=True)
  (relu2): ReLU()
  (fc3): Linear(in_features=10, out_features=1, bias=True)
  (sigmoid): Sigmoid()
)
```

```
from torch import nn
from collections import OrderedDict

# define model architecture
model3 = nn.Sequential(OrderedDict([
    ('fc1', nn.Linear(16, 12)),
    ('relu1', nn.ReLU()),
    ('fc2', nn.Linear(12, 10)),
    ('relu2', nn.ReLU()),
    ('fc3', nn.Linear(10, 1)),
    ('sigmoid', nn.Sigmoid())
]))

# print model architecture
print(model3)
```

```
Sequential(
  (fc1): Linear(in_features=16, out_features=12, bias=True)
  (relu1): ReLU()
  (fc2): Linear(in_features=12, out_features=10, bias=True)
  (relu2): ReLU()
  (fc3): Linear(in_features=10, out_features=1, bias=True)
  (sigmoid): Sigmoid()
)
```

```
import torch
import torch.nn.functional as F
from torch import nn

# define the network class
class MyNetwork(nn.Module):
    def __init__(self):
        # call constructor from superclass
        super().__init__()

        # define network layers
        self.fc1 = nn.Linear(16, 12)
        self.fc2 = nn.Linear(12, 10)
        self.fc3 = nn.Linear(10, 1)

    def forward(self, x):
        # define forward pass
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = torch.sigmoid(self.fc3(x))
        return x

    # instantiate the model
model1 = MyNetwork()

# print(model1)

MyNetwork(
  (fc1): Linear(in_features=16, out_features=12, bias=True)
  (fc2): Linear(in_features=12, out_features=10, bias=True)
  (fc3): Linear(in_features=10, out_features=1, bias=True)
)
```

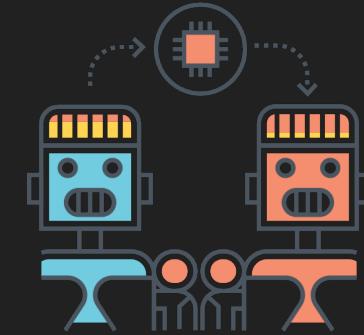
Types of Learning



Supervised
Learning



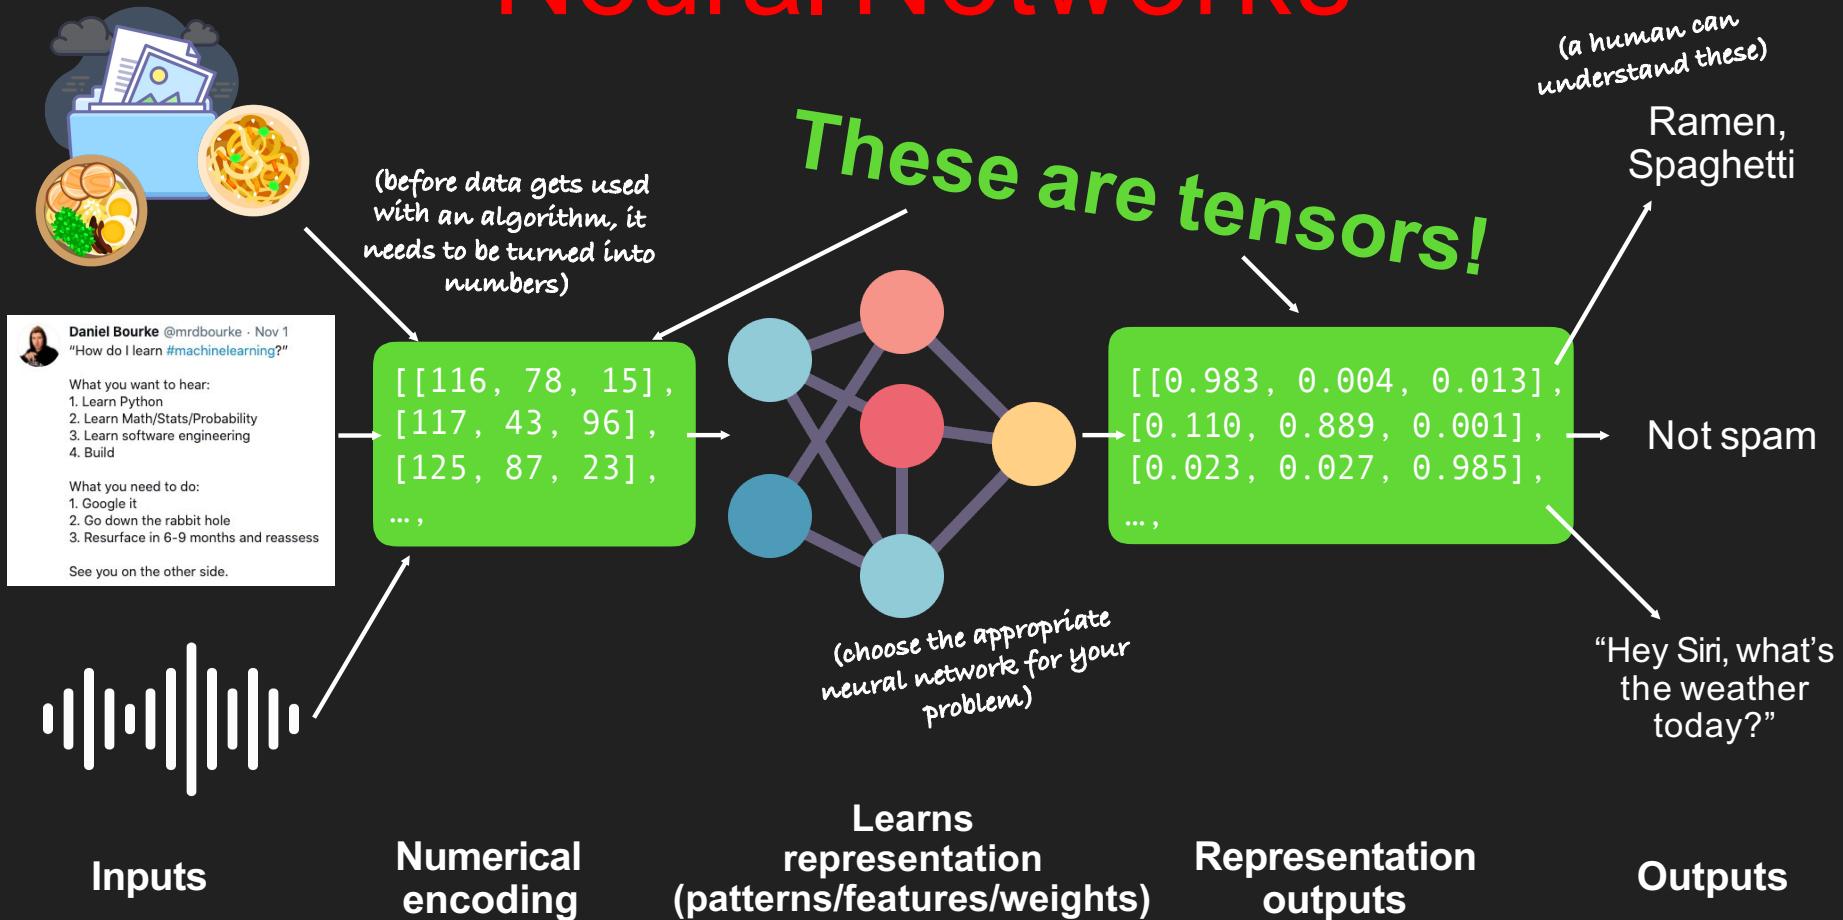
Unsupervised &
Self-supervised
Learning

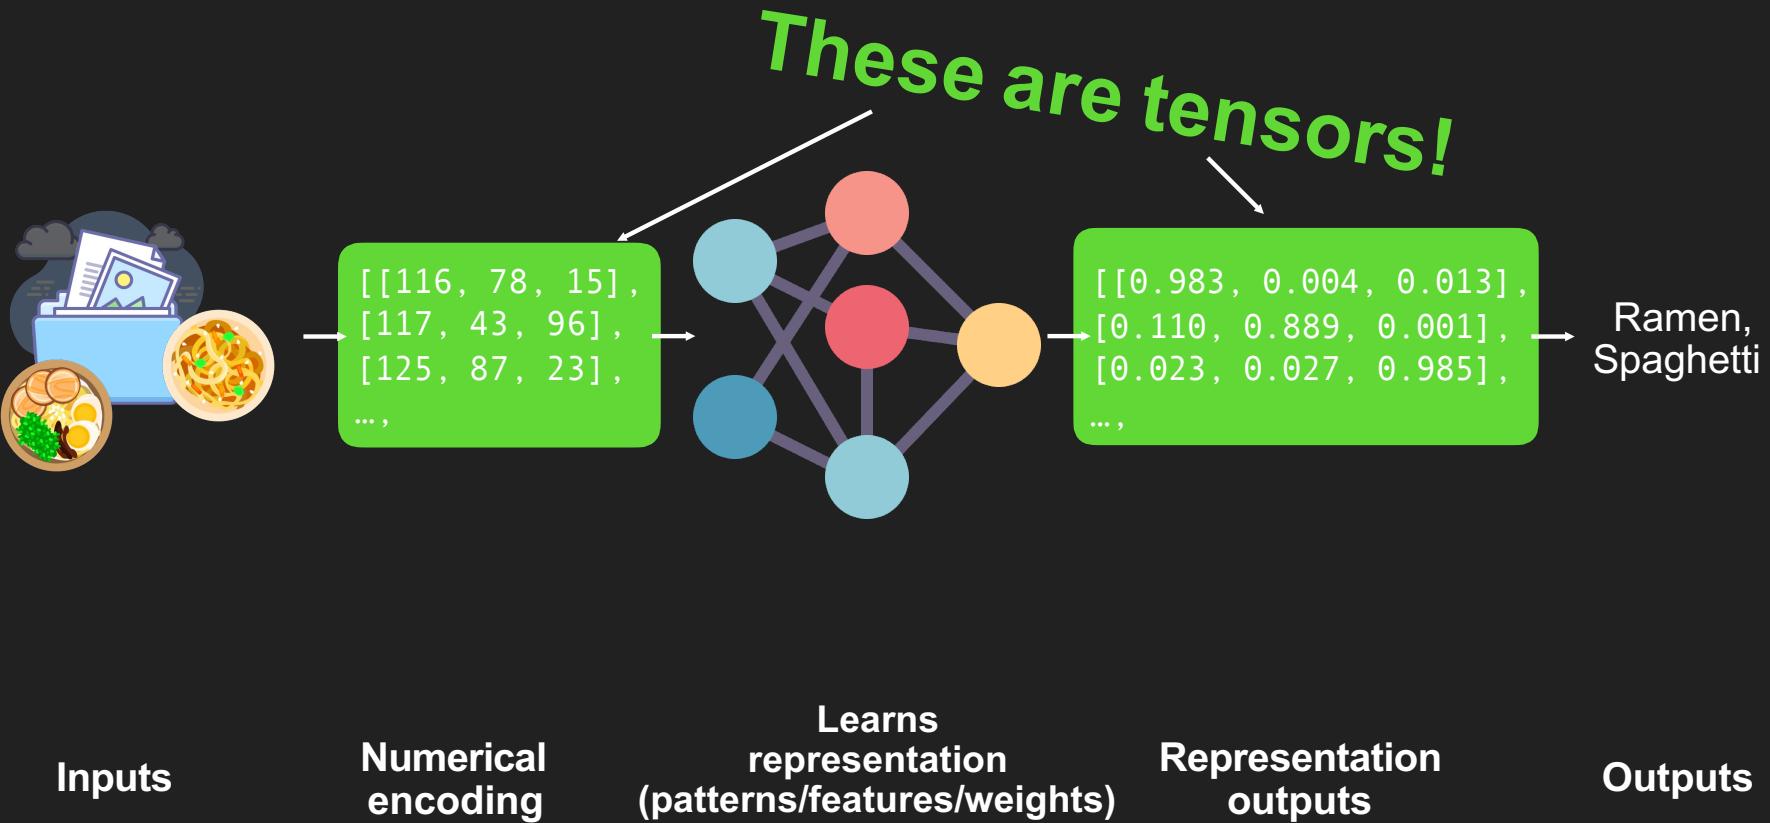


Transfer
Learning

We'll be writing code to do these,
but the style of code can be adopted across learning paradigms.

Neural Networks

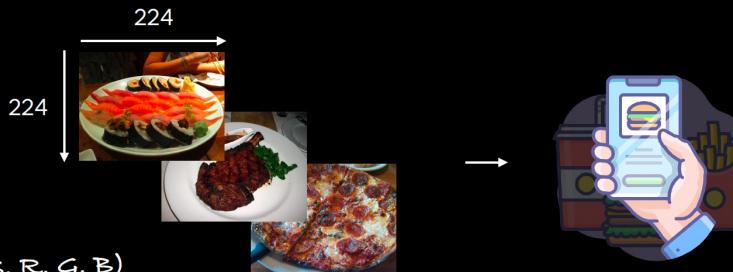




Classification inputs and outputs

$W = 224$
 $H = 224$
 $C = 3$

(c = colour channels, R, G, B)

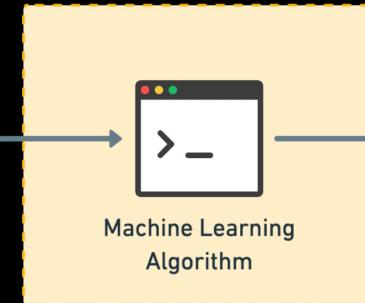


Sushi 🍣
Steak 🥩
Pizza 🍕

Actual output

$[[0.31, 0.62, 0.44...],$
 $[0.92, 0.03, 0.27...], \rightarrow$
 $[0.25, 0.78, 0.07...],$
 $\dots,$ (normalized pixel values)

Numerical encoding



(often already exists, if not,
you can build one)

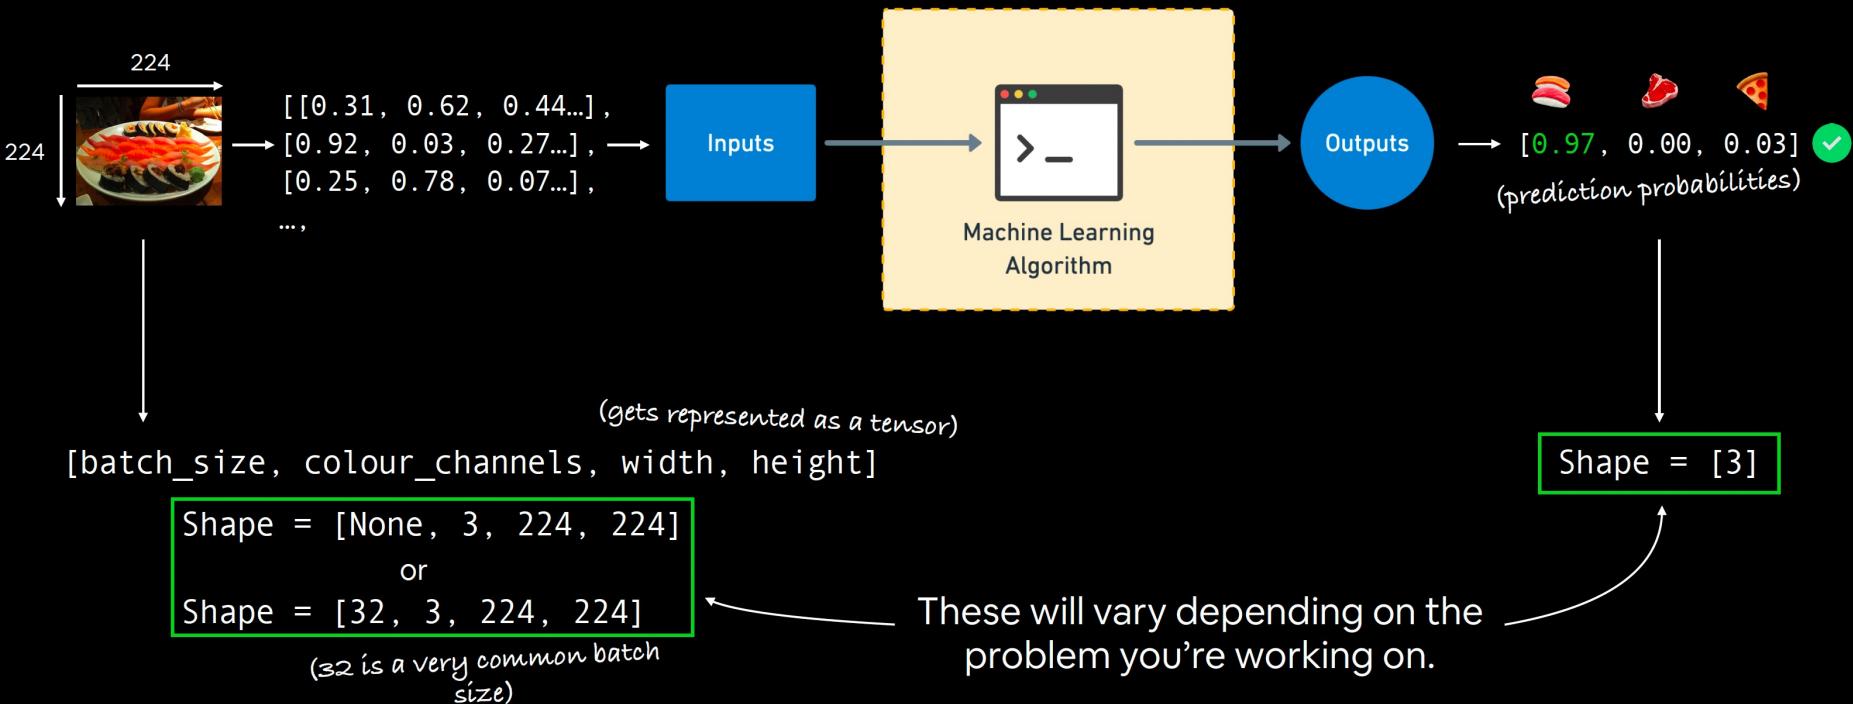
[[0.97, 0.00, 0.03], ✓
[0.81, 0.14, 0.05], ✗
[0.03, 0.07, 0.90], ✓
 $\dots,$

Predicted output

(comes from looking at lots
of these)

Input and output shapes

(for an image classification example)



Example: Regression

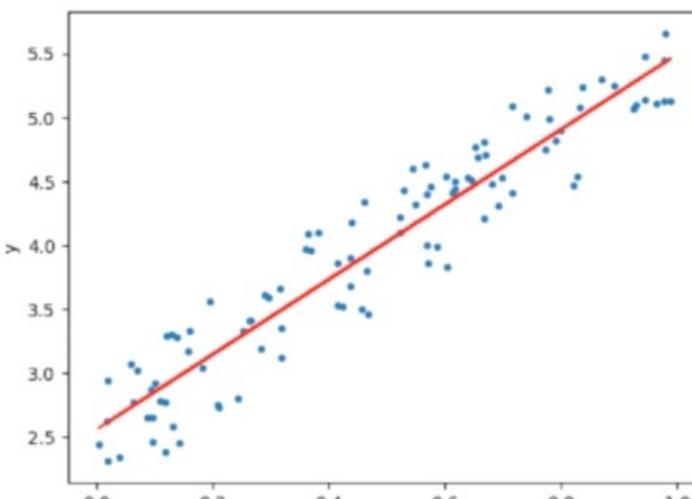
- Fit a line or curve
- This is why grumpy old statisticians like to say: “machine learning is just glorified curve-fitting”

$$\hat{y} = mx + b$$

(sometimes we also use “a” for the slope)

Salary

Years of experience

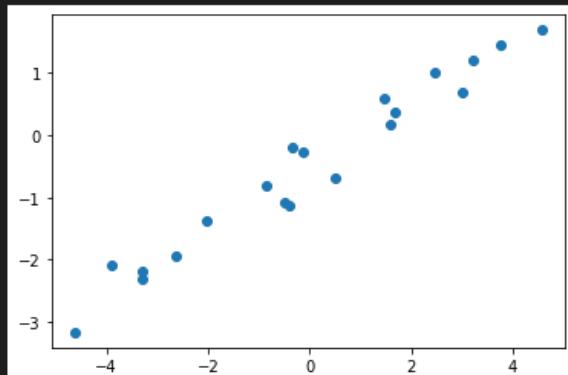


```
import matplotlib.pyplot as plt  
import numpy as np  
✓ 0.4s
```

```
# generate 20 data points  
N = 20  
  
#random data on x-axis  
x= np.random.rand(N)*10-5  
  
y = 0.5*x -1+ np.random.rand(N)  
✓ 0.2s
```

```
plt.scatter(x,y)  
✓ 0.1s
```

```
<matplotlib.collections.PathCollection at 0x7fc1b147a410>
```



```
# In ML we want our data to be of shape:  
# (num_samples x num_dimensions)  
X = x.reshape(N, 1)  
Y = y.reshape(N, 1)
```

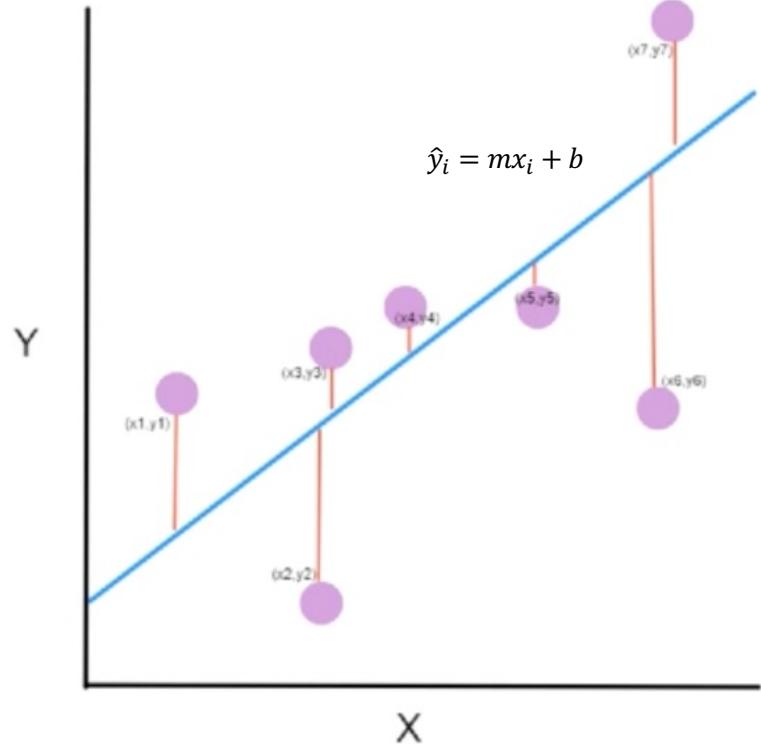
```
✓ 0.2s
```

The Loss

- Data = $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- N = number of samples in the dataset
- The line **cannot** perfectly pass through all the data points

MSE = Mean Squared Error

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$



Applying the MSE to find slope / intercept

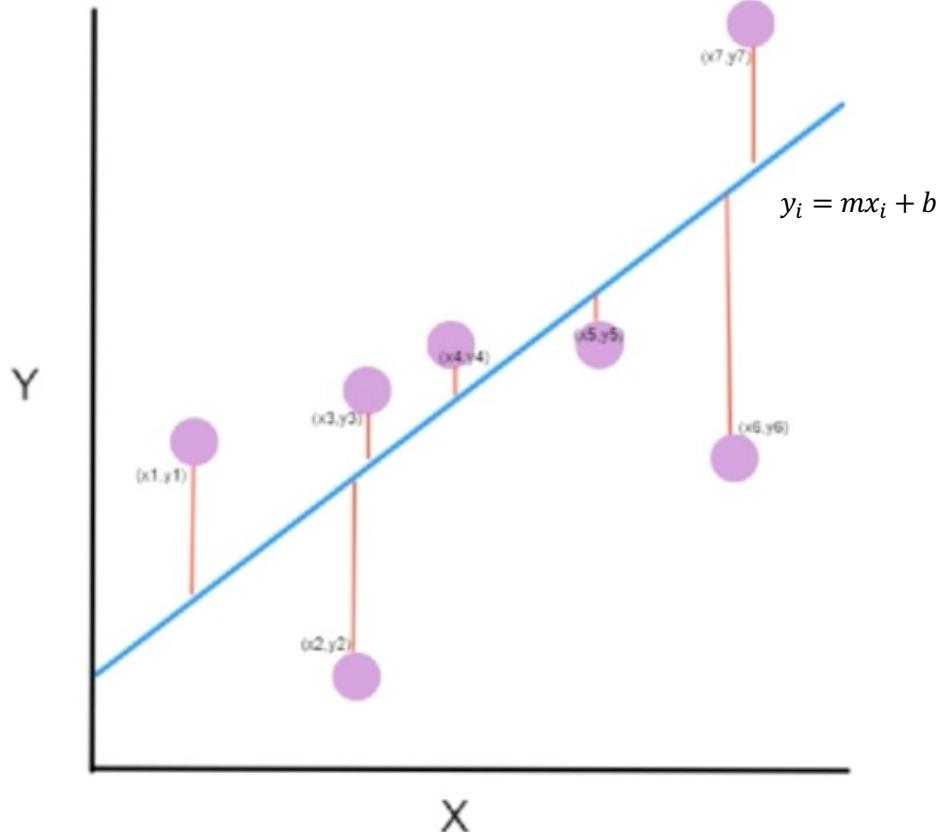
- Plug in the expression for the predictions (\hat{y}_i)
- Quiz: What are the **variables** in this expression?

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

Applying the MSE to find slope / intercept

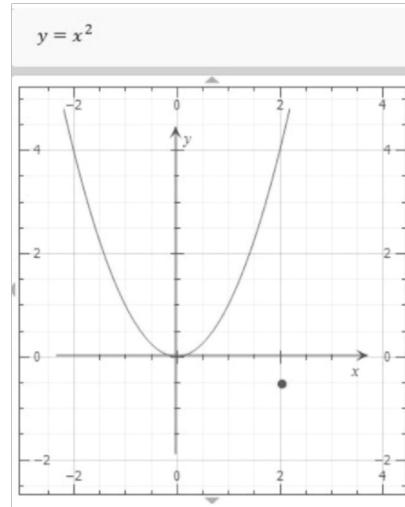
- Note: L (loss) = MSE
- We want to *minimize the loss* with respect to the parameters (m, b)

$$m^*, b^* = \arg \min_{m, b} L$$



Minimize $f(x) = x^2$

First Solution



- Use calculus!
- $df / dx = 2x = 0$
- Solve for x : $x = 0$

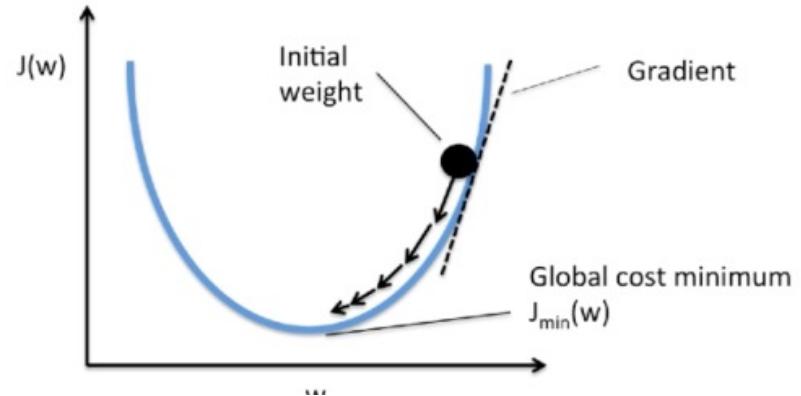
Second Solution : Gradient Descent

- Find the derivatives and set them to 0, solve for the parameters
- Yields 2 equations and 2 unknowns: can solve for (m, b)
- Try it as an exercise!: Should be in terms of $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$

$$\partial L/\partial m = 0, \partial L/\partial b = 0$$

Gradient Descent to find hidden parameters

```
# gradient descent  
# pseudocode  
  
 $\theta = (m, b) = \text{random}()$   
for i in range(n_epochs):  
     $\theta = \theta - \eta \nabla_{\theta} L$ 
```



$$L = MSE = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

$$\frac{\partial}{\partial m} = \frac{2}{n} \sum_{i=1}^n -x_i(y_i - (mx_i + b))$$

$$\frac{\partial}{\partial b} = \frac{2}{n} \sum_{i=1}^n -(y_i - (mx_i + b))$$

```
def Gradient(m,b) :  
    grad_m =0; grad_b = 0;  
  
    for i in range(N) :  
        grad_m += -X[i]*(Y[i]- (m*X[i]+b))  
        grad_b += - (Y[i]- (m*X[i]+b) )  
  
    return grad_m, grad_b  
  
Gradient(-1,-1)  
✓ 0.3s  
(array([-209.28382035]), array([-10.12894605]))
```

```
# gradient descent
# pseudocode

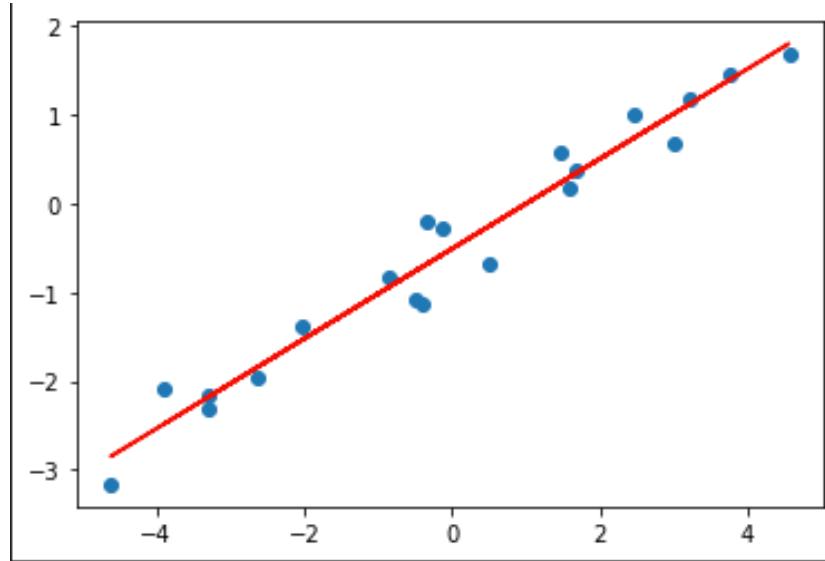
θ = (m, b) = random()
for i in range(n_epochs):
    θ = θ - η∇_θ L
```

```
n_epochs      = 1000
learning_rate = 0.01
theta         = np.random.rand(2,1)

for i in range(n_epochs) :
    grad_m, grad_b = Gradient(theta[0], theta[1])
    theta[0] -= learning_rate*grad_m
    theta[1] -= learning_rate*grad_b

theta
✓ 0.1s
array([[ 0.50714187],
       [-0.51174366]])
```

```
predict = theta[0]*X + theta[1]
plt.plot(X, predict, 'r')
plt.scatter(x,y)
✓ 0.8s
```



Third Solution :Pytorch

$$\theta = \theta - \eta \nabla_{\theta} \mathcal{L}$$

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

```
# Loss and optimizer
criterion = nn.MSELoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

```
# gradient descent
# pseudocode

θ = (m, b) = random()
for i in range(n_epochs):
    θ = θ - η∇_θ L
```

```
# Train the model
n_epochs = 30
for it in range(n_epochs):
    # zero the parameter gradients
    optimizer.zero_grad()

    # Forward pass       $\hat{y}_i = mx_i + b$ 
    outputs = model(inputs)
    loss = criterion(outputs, targets)

    # Backward and optimize
    loss.backward()
    optimizer.step()
```

```

1 # Train the model
2 n_epochs = 30
3 losses = []
4 for it in range(n_epochs):
5     # zero the parameter gradients
6     optimizer.zero_grad()
7
8     # Forward pass
9     outputs = model(inputs)
10    loss = criterion(outputs, targets)
11
12    # keep the loss so we can plot it later
13    losses.append(loss.item())
14
15    # Backward and optimize
16    loss.backward()
17    optimizer.step()
18
19    print(f'Epoch {it+1}/{n_epochs}, Loss: {loss.item():.4f}')

```

Epoch 1/30, Loss: 12.1352
Epoch 2/30, Loss: 2.8066
Epoch 3/30, Loss: 2.4554
Epoch 4/30, Loss: 2.1753
Epoch 5/30, Loss: 1.9483
Epoch 6/30, Loss: 1.7641
Epoch 7/30, Loss: 1.6148

```

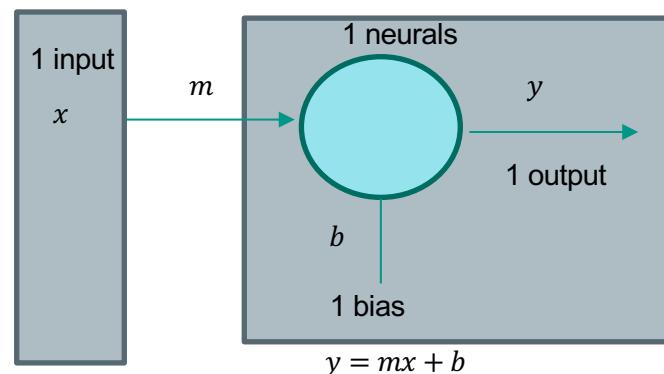
[ ] 1 # Create the linear regression model
2 model = nn.Linear(1, 1)

[ ] 1 # Loss and optimizer
2 criterion = nn.MSELoss()
3 optimizer = torch.optim.SGD(model.parameters(), lr=0.05)

[ ] 1 # In ML we want our data to be of shape:
2 # (num_samples x num_dimensions)
3 X = X.reshape(N, 1)
4 Y = Y.reshape(N, 1)
5
6 # PyTorch uses float32 by default
7 # Numpy creates float64 by default
8 inputs = torch.from_numpy(X.astype(np.float32))
9 targets = torch.from_numpy(Y.astype(np.float32))

net = torch.Linear(1,1)

```



Example : Parameter Estimation

Home Work 2

$$\text{signal} = A \sin(t + \theta) + b$$

$$y = \text{signal} + \text{noise}$$

```
t = np.arange(0, 20, 0.1)

#random data on x-axis
A = np.random.rand()
b = 20*np.random.rand()
ceta = 2*np.random.rand()*np.pi

signal = A*np.sin(t+ceta)+b
y = signal + 0.1*np.random.randn(len(t))

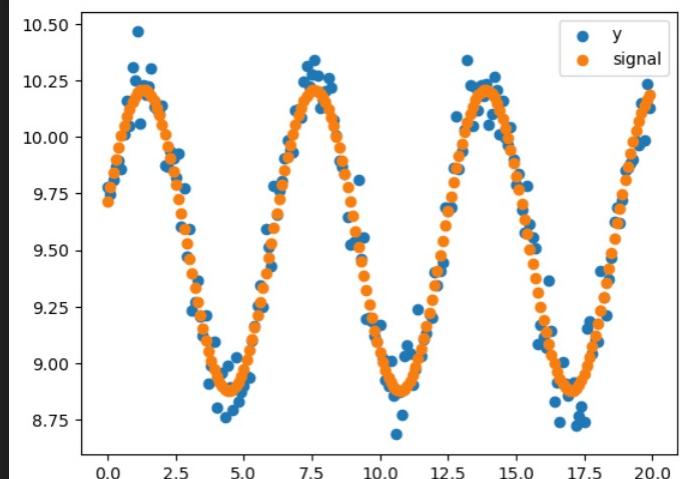
print(f"A = {A}, ceta = {ceta} b= {b}")
```

A = 0.5929415692077172, ceta = 5.396419680771073 b= 6.038952373195254

```
plt.scatter(t,y)
plt.scatter(t,signal)
plt.legend(['y','signal'])
```

✓ 0.1s

<matplotlib.legend.Legend at 0x16aaaf3100>

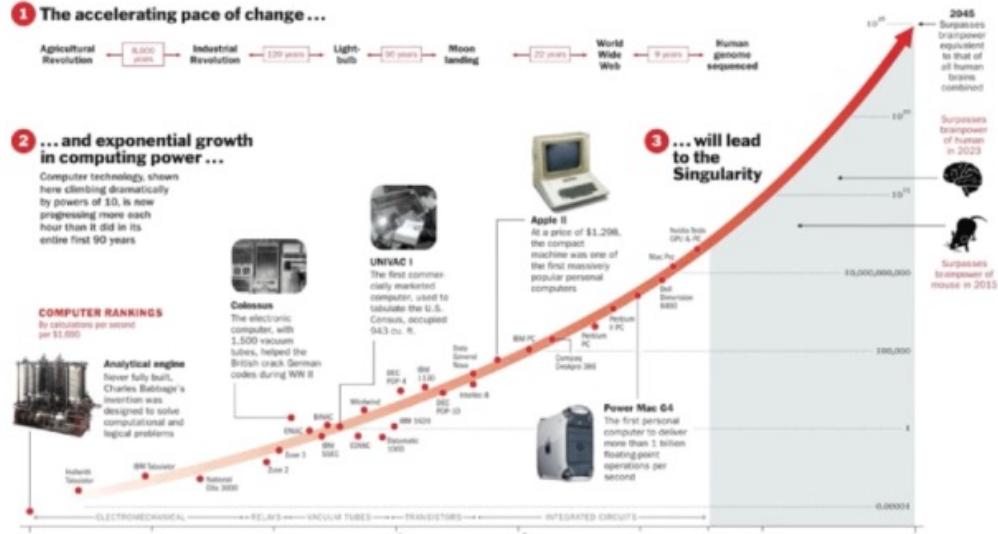




Home Work

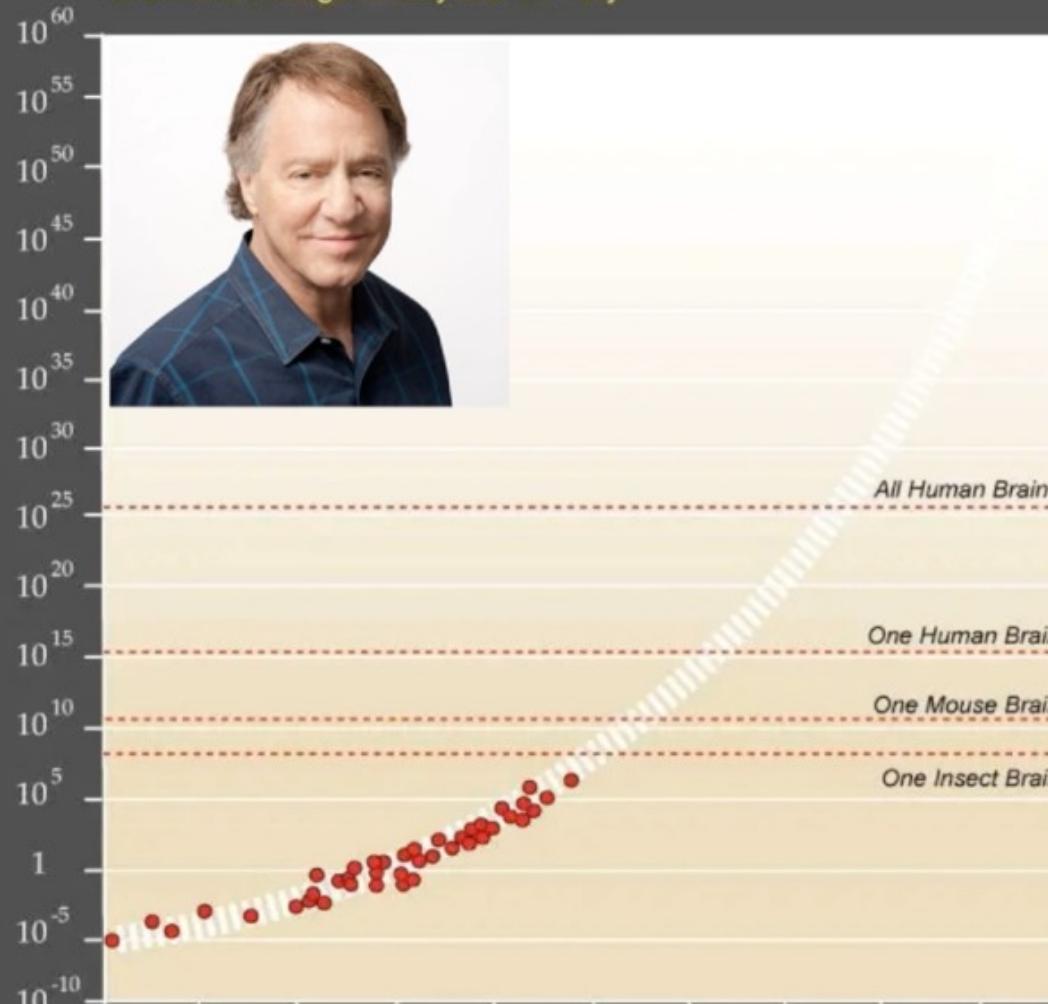
Real-World Dataset: Moore's Law

- The number of transistors per square inch on integrated circuits doubles approximately every 2 years

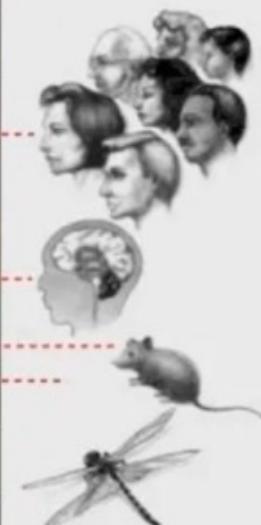


Exponential Growth of Computing

Twentieth through twenty first century



Logarithmic Plot



Home Work 1

```
df = pd.read_csv('https://raw.githubusercontent.com/lazyprogrammer/machine_learning_examples/master/tf2.0/moore.csv', header=None)
df= df.rename(columns={0: "years", 1: "data"})
```

```
df
```

```
✓ 0.2s
```

	years	data
0	1971	2300
1	1972	3500
2	1973	2500
3	1973	2500
4	1974	4100
...
157	2017	18000000000
158	2017	19200000000
159	2018	8876000000
160	2018	23600000000
161	2018	9000000000

```
162 rows x 2 columns
```

```
df.plot(kind='scatter', x='years', y='data')
```

```
✓ 0.9s
```

```
<AxesSubplot:xlabel='years', ylabel='data'>
```

